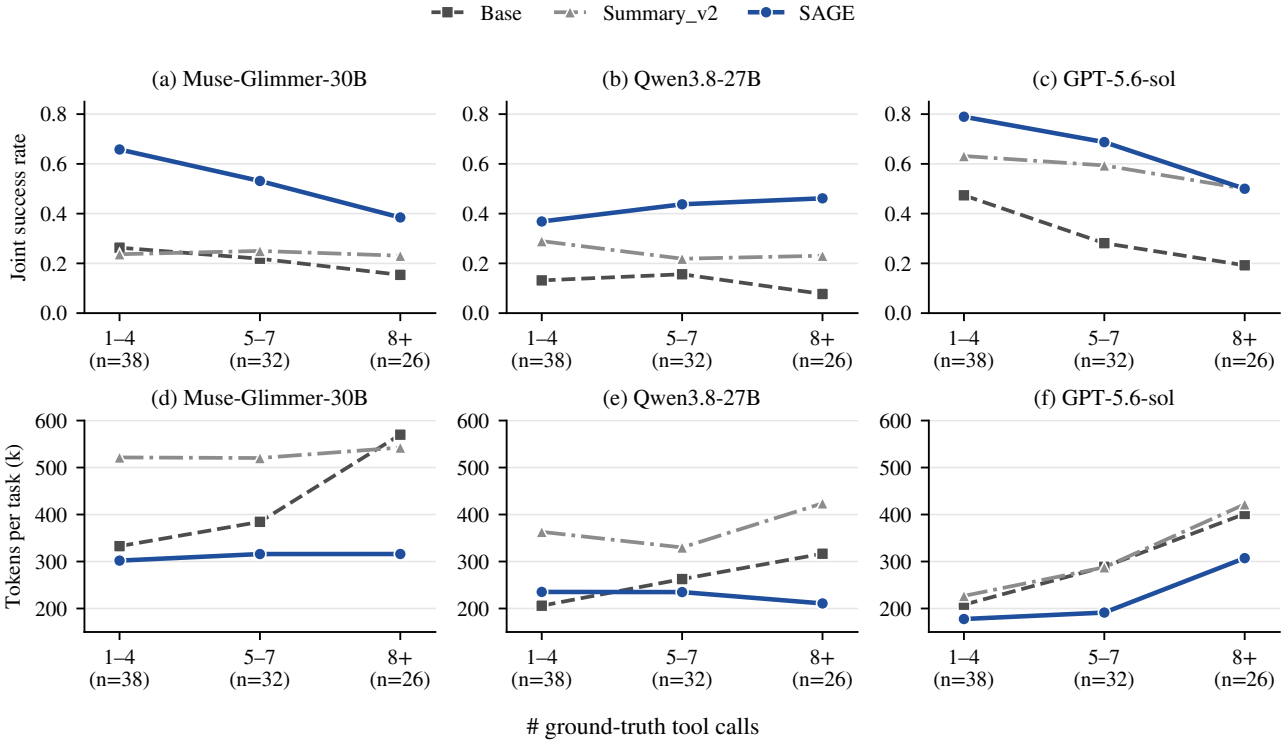
Base
Summary_v2
SAGE
(a) Muse-Glimmer-30B
(b) Qwen3.8-27B
(c) GPT-5.6-sol
(d) Muse-Glimmer-30B
(e) Qwen3.8-27B
(f) GPT-5.6-sol
Joint success rate
Tokens per task (k)
0.0
0.2
0.4
0.6
0.8
200
300
400
500
600
1–4 (n=38)
5–7 (n=32)
8+ (n=26)
ground-truth tool calls